KUSH JAIN

kdjain@andrew.cmu.edu • (650) 965 1750 • https://www.linkedin.com/in/kush-jain/• https://www.kushjain.com

RESEARCH INTERESTS: My research lies at the intersection of natural language processing, software engineering and testing. Specific topics include machine learning for code and tests, mutation testing, fuzzing.

EDUCATION

Carnegie Mellon University – School of Computer Science   August 2021 – May 2026 (expected)
Doctor of Philosophy, Computer Science, advisor: Claire Le Goues
University of Texas at Austin   August 2018 – May 2021
Batchelor of Science, Computer Science, advisor: Milos Gligoric

PUBLICATIONS

[1] CAT-LM: Training Language Models on Aligned Code And Tests
38th IEEE/ACM International Conference on Automated Software Engineering (ASE 2023)
Nikitha Rao*, **Kush Jain***, Uri Alon, Claire Le Goues, Vincent Hellendoorn (* = equal contribution)

[2] Contextual Predictive Mutation Testing
31st ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (FSE 2023)
Kush Jain, Uri Alon, Alex Groce, and Claire Le Goues

[3] Mind the Gap: The Difference Between Coverage and Mutation Score Can Guide Testing Efforts
34th International Symposium on Software Reliability Engineering (ISSRE 2023)
Kush Jain, Goutamkumar Tulajappa Kalburgi, Claire Le Goues, Alex Groce

[4] Looking for Lacunae in Bitcoin Core's Fuzzing Efforts
43rd IEEE/ACM International Conference on Software Engineering (ICSE 2022)
Alex Groce, **Kush Jain**, Rijnard van Tonder, Goutamkumar Tulajappa Kalburgi, and Claire Le Goues

[5] Registered Report: First, Fuzz the Mutants
1st International Fuzzing Workshop (FUZZING 2022)
Alex Groce, Goutamkumar Tulajappa Kalburgi, Claire Le Goues, **Kush Jain**, and Rahul Gopinath

[6] Programming and Execution Models for Parallel Bounded Exhaustive Testing
Conference on Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA 2021)
Nader Al Awar, **Kush Jain**, Christopher J. Rossbach, and Milos Gligoric

[7] mCoq: Mutation Analysis for Coq Verification Projects
42nd IEEE/ACM International Conference on Software Engineering Tool Demonstrations Track (ICSE 2021)
Kush Jain, Karl Palmskog, Ahmet Celik, Emilio Jesus Gallego Arias, and Milos Gligoric

EXPERIENCE

IBM TJ Watson – *AI Research Intern;* Yorktown Heights, New Yrok   June 2023 – August 2023
- Developed a novel LLM retrieval prompting approach that produces correct and diverse OpenAPI parameter examples
- Improved state of the art in a wide range of domains including fuzzing, chatbot dialog systems and human API understanding

Amazon Lab126 – *Software Engineering Intern;* Sunnyvale, California   June 2021 – August 2021
- Developed a webapp to manage the approval process for all prototype devices at Amazon
- Migrated data to DynamoDB and integrated unified authentication

Amazon Lab126 – *Software Engineering Intern;* Sunnyvale, California June 2020 – August 2020
- Developed a device search service for prototype devices using AWS lambda, ElasticSearch, API gateway and Database Migration Service to serve over 30 million requests a month, while dramatically improving existing search functionality in a schema change tolerant way, leveraging federated authentication

VISA Inc. – *Software Engineering Intern;* Austin, Texas June 2019 – August 2019
- Developed a dashboard to track health of core IT services using NodeJS, React and PowerShell. In the first two months of production, proactively detected five major outages, preventing over 250 support tickets
- Implemented a customized link shortener for VISA's internal network, using NodeJS, React and SQL.

OpsHub Inc. – *Software Engineering Intern;* Palo Alto, California June 2018 – August 2018
- Prototyped an Angular dashboard to visualize multi-system KPI's using the company's integration platform
- Proof of concept was successful, and company is looking to fully develop the product

OpsHub Inc. – *Software Engineering Intern;* Palo Alto, California June 2017 – August 2017
- Developed a model to analyze the riskiness of a source code file and to predict the number of bugs expected
- Got 70% accuracy and had visibility to improve it further by bringing in data from additional systems

OPEN SOURCE CONTRIBUTIONS

FuzzBench: added our fuzzers that fuzz mutants first and then run normal fuzzing over the benchmarks – https://github.com/google/fuzzbench

mCoq: mutation analysis tool for Coq verification projects, used by around a dozen developers from around the world – https://github.com/EngineeringSoftware/mcoq

SMUM-Checkin: open source code of our digital recordkeeping system used by Santa Maria Urban Ministry – https://github.com/UnconditionedLife/smum

SKILLS
- Expert in Java, Python (including ML libraries: pandas, sklearn, numpy, pytorch), C, and C++
- Proficient in Angular, CSS, HTML, Javascript, JQuery, Selenium, SQL, React, Ruby

PROJECTS

[1] CAT-LM – Training Language Models on Aligned Code and Tests (ASE 2023) August 2022 – May 2023
Collaborators: Vincent Hellendoorn (CMU), Claire Le Goues (CMU)
- Testing is an integral part of software development, yet has received far less attention than areas such as code completion and bug repair
- We demonstrate the importance of incorporating domain knowledge, in this case the relationship between code and test files to create more powerful models for automated test generation
- CAT-LM outperforms state of the art test generation models on both lexical and runtime metrics

[2] Predictive Mutation Testing Through Large Language Models (FSE 2023) May 2022 – May 2023
Collaborators: Uri Alon (CMU), Claire Le Goues (CMU)
- Mutation testing has been extensively researched in academia, yet has failed to achieve much industry traction due to its large compute requirements
- Our idea is to leverage large language models over code to embed methods, operators, and test suites to predict whether a given mutant will be killed or not killed without running the test suite, saving significant compute
- Our novel approach outperforms state of the art techniques in this space, with major gains in time savings and model accuracy

[3] Analyzing the Difference Between Code Coverage and Mutation Score (ISSRE 2023) August 2021 – August 2022
Collaborators: Claire Le Goues (CMU), Alex Groce (NAU)

- Most testing efforts in industry use coverage as the primary metric of test-adequacy, with an even smaller subset using mutation score, however neither metric is fully sufficient
- We analyze the difference between code coverage and mutation score, and examine how it can guide real world testing efforts
- We find that this difference can detect both induced examples of testing inadequacy and real-world examples corresponding to GitHub commits

[4] Looking for Lacunae in Bitcoin Core's Fuzzing Efforts (ICSE 2022)   August 2021 – August 2022
Collaborators: Claire Le Goues (CMU), Alex Groce (NAU)
- Conducted cross-cryptocurrency mutation testing, benchmarking Bitcoin against other cryptocurrencies to understand its relative strengths, which underlined Bitcoin Core's superior file and project coverage
- Manually examined cases where coverage and mutation score diverged, pinpointing gaps in cryptocurrency testing efforts

[5] First, Fuzz the Mutants (FUZZING 2022)   August 2021 – May 2022
Collaborators: Claire Le Goues (CMU), Alex Groce (NAU)
- Fuzzing is a popular testing technique that has begun to achieve more widespread adoption, however fuzzers still struggle to reach certain code snippets due to number or branches or expensive compile optimizations
- Our approach is to first apply mutation analysis to the target program being fuzzed, and then fuzz these mutants for half of the fuzzing budget to generate seeds that cover these hard-to-reach branches
- Results indicate that this approach can both lead to higher coverage and more bugs found on multiple fuzzing benchmarks

[6] Tempo – Bounded Exhaustive Testing (OOPSLA 2021)   August 2020 – August 2021
Collaborators: Christopher J. Rossbach (UT Austin), Milos Gligoric (UT Austin)
- We develop Tempo: a programming and execution model for parallel bounded exhaustive testing. It supports hybrid test generation programs using two strategies: fork-based and re-execution based
- Evaluation shows fork-based works well for simple programs on GPUs while re-execution based handles complex programs better
- We also study Clang and GCC finding multiple bugs in both compilers

[7] Mutation Analysis for Coq Verification Projects (ICSE 2021)   August 2020 – August 2021
Collaborators: Emilio Jesus Gallego Arias (IRIF), Milos Gligoric (UT Austin)
- We perform mutation analysis on Coq formal specifications to detect partial or incomplete specifications
- Implemented a wide range of mutation operators ranging from simple (replace and with or) to complex ones (reversing Coq inductive cases)